



solutions_V
2

2022 Data Systems Final Exam (Solutions)

Jaylitha C¹

November 22, 2022

Contents

Question 1 2

Question 2 4

Question 3 7

Question 4 9

¹jaylitha.cs@research.iiit.ac.in

Question 1

Let R_1, R_2, \dots, R_k be k relations, each having attribute A . Let A have domain values $dom(A) = \{d_1, d_2, \dots, d_p\}$.

Let $(n_{1j}, n_{2j}, \dots, n_{kj})$ be the number of rows, respectively, for domain values $\{d_1, d_2, \dots, d_p\}$ of attribute A in relation R_i . That is, n_{ij} is the number of rows in R_i when attribute A takes domain value d_j . Let

$$T = R_1 \bowtie_{R_1.A=R_2.A} R_2 \bowtie_{R_2.A=R_3.A} R_3 \dots R_{k-1} \bowtie_{R_{k-1}.A=R_k.A} R_k$$

Compute the maximum number and the minimum number of rows for T . State which strategy (if any) can be used to efficiently compute T .

Answer. The number of rows in T is given by the following equation

$$\sum_{j=1}^p \prod_{i=1}^k n_{ij} \quad (1)$$

For each value d_j in the domain (A), let there are $n_{1j} \times n_{2j} \times \dots \times n_{kj}$ tuples. Eq (1) is the minimum and maximum number of tuples in T without imposing any constraints on the "histograms" of the relations. Assuming we can, then the minimum number of tuples of T is 0. This minimum can occur in the extreme case where one relation is empty. However, the requirement for this minimum case is stated below

$$\forall d_j \in dom(A) \exists R_i, n_{ij} = 0$$

For each value d_j in the domain, there should exist at least one relation that does not contain any tuples with value d_j for A . The maximum number of tuples of T in theory is the cross product of all the relations. Let N_1, N_2, \dots, N_k be the number of tuples of each relation. Let $n_{1j} = N_1$ and $n_{ij} = 0$ for all $j \neq 1$ i.e. every tuple of every relation has value d_1 for attribute A . Then the number of tuples of T is

$$\prod_{i=1}^k N_i \quad (2)$$

Which strategy would I use to efficiently compute T ? The objective is to reduce the number of block accesses. To do this, we can split the problem up into disjoint join operations. We first partition ALL the relations using the same partitioning function. The partitioning

Eq (1) can also be rewritten in terms of the histogram values as so

$$\sum_{j=1}^p \prod_{i=1}^k h_{ij}$$

This form relates to Eq (1) using the follow property: Let $d_1, d_2, \dots, d_{n_1}, d_{n_1+1}, \dots, d_{n_1+n_2}, \dots$ be non-negative numbers, then

$$\begin{aligned} n_1 h_{1j} + n_2 h_{2j} &\leq n_1 h_{1j} + n_2 h_{2j} + n_3 h_{3j} + \dots + n_k h_{kj} \\ &= (n_1 + n_2) h_{2j} + n_3 h_{3j} + \dots + n_k h_{kj} \end{aligned}$$

2022 DATA SYSTEMS FINAL EXAM (SOLUTIONS) 3

Figure 1: Partitioned multi-join example

function maps the each value in the domain of A to one partition (for instance a hash or block/range partition). Take for example the three $k=3$ relations shown in Fig 1. The domain of attribute A is $\{1, 2, 3\}$. Every relation is partitioned using the same partition function $h(A) = A \bmod 3$. Relation R_1 generates three partitions R_1^1 of tuples that hash to 1, R_1^2 of tuples that hash to 2 and R_1^3 of tuples that hash to 0. These partitions are written back to the disk. R_2 and R_3 undergo the same partitioning. Fig. 1 shows the results of partitioning the relations. It is important to note that tuples that correspond to one partition do not join with tuples in any other partition. For instance, tuples of R_1^1 only join with tuples of R_2^1 and R_3^1 . Therefore, the large multi-join has been split up into smaller multi-joins (or in this case, cartesian product) operations of the partitions.

if any partition for any relation is empty, the result is empty.

¹For instance, if a partition is too big, we can check to see if smaller partitions can be created. If not, the implementation of the operator needs to accommodate this case. If the partitions are too small, then too many block accesses will take place and we will not utilize the main memory efficiently. To handle this, smaller partitions can be combined using some bin-packing algorithm. The histogram will aid this optimization.

2022 DATA SYSTEMS FINAL EXAM (SOLUTIONS) 4

Question 2

Question 2. Consider a relation R having 10^6 ($1M$) rows with a non-null attribute B having domain values $\{YES, NO\}$. The data page size is 2000, row length for R is 100. The length of attribute B is 10, and the length of the record pointer and block pointer is 10 bytes each. Let there be a single-level secondary index on B . The amount of buffer space for storing the index is much smaller than the size of the index. What are the minimum and maximum number of data page accesses needed to execute

SELECT * FROM R WHERE B = "YES";

using the single level secondary index (NOT B+ tree) on B , when there are 9,99,980 rows with the value "YES", when there are 40,000 rows with the value "YES", and when there is one row with the value "YES".

Answer. In theory, the secondary index can be represented as a set of $\langle K, P \rangle$ pairs where K is the key ($K \in \{YES, NO\}$) and P is a record or block pointer that point to a tuple with value K . Using the implementation specified in the textbook¹ assume indirection blocks are used to store the pointers. Each index entry points to a block which contains block or record pointers². If the number of pointers are large, then the index points to a linked list of indirection blocks that hold pointers.

Since we are optimizing for the number of blocks accessed, storing record pointers over block pointers increases the size of the index and the number of block accesses. In the best case, block pointers are used and in the worst case, record pointers are used.

The total number of block accesses needed to execute the query is the sum of the number of index blocks and R data blocks accessed.

INDEX BLOCKS + # DATA BLOCKS

Some math:

The number of data records per page = $\left\lfloor \frac{2000}{100} \right\rfloor = 20$

The number of pages needed to store $R = \left\lceil \frac{10^6}{20} \right\rceil = 5 \times 10^4$

The number of block or record pointers per page = $\left\lceil \frac{2000}{10} \right\rceil = 200$

The number of index records per page ($\langle K, P \rangle$) = $\left\lceil \frac{2000}{10} \right\rceil = 100$

¹E Elmasri and Shamkant B Navathe. *Fundamentals of Database Systems*. Springer, 2000. Accessed on 19/11/2022. Link.

²According to the textbook, a record pointer contains a block pointer.

If record pointers are used, then the semantics associated with record pointers are followed. To elaborate, when executing the query, for each record pointer, the relevant block is read into memory. Therefore, the same block can be read into main memory multiple times if multiple record pointers point to the same block.

Therefore, the minimum number of block accesses is = $5 \times 10^4 - 1 + 20 = 50290$

In the worst case, all blocks are accessed and a few are accessed multiple times. Specifically, the worst case needs 999980 block accesses. Since record pointers are used, 999980 record pointers are stored in the index. The number of index blocks accessed is

$$\left\lceil \frac{999980}{200} \right\rceil = 1 + \lceil 4999.9 \rceil = 5001$$

Therefore, the maximum number of block accesses is = $999980 + 5000 = 1004980$

YES1 = 40000

Again, in the best case these records are stored contiguously. The number of data block accesses is

$$\left\lceil \frac{40000}{20} \right\rceil = 2000$$

and the index holds block pointers for each of these 2000 blocks. The number of index block accesses is

$$1 + \left\lceil \frac{2000}{200} \right\rceil = 11$$

Therefore, the minimum number of block accesses is = $2000 + 11 = 2011$

In the worst case, we access at most 40000 blocks³. 4K block pointers

³We cannot partition 40000 rows into 50000 blocks such that each block gets at least one row

2022 DATA SYSTEMS FINAL EXAM (SOLUTIONS) 5

There are only 2 index records, therefore, only one block is needed to store the index records (more blocks are needed to store the indirection blocks)

YES1 = 999980

In the best case, all YES records are stored contiguously in $5 \times 10^4 - 1$ blocks. The number of indirection blocks accessed is

$$\left\lceil \frac{5 \times 10^4 - 1}{200} \right\rceil = \left\lceil \frac{5 \times 10^4}{200} - \frac{1}{200} \right\rceil = \left\lceil 250 - \frac{1}{200} \right\rceil = 250$$

The number of index blocks accessed is = $1 + 250 = 251$. The additional block access is necessary to bring the index records into memory.

Therefore, the minimum number of block accesses is = $5 \times 10^4 - 1 + 250 = 50290$

In the worst case, all blocks are accessed and a few are accessed multiple times. Specifically, the worst case needs 999980 block accesses. Since record pointers are used, 999980 record pointers are stored in the index. The number of index blocks accessed is

$$\left\lceil \frac{999980}{200} \right\rceil = 1 + \lceil 4999.9 \rceil = 5001$$

Therefore, the maximum number of block accesses is = $999980 + 5000 = 1004980$

YES1 = 40000

Again, in the best case these records are stored contiguously. The number of data block accesses is

$$\left\lceil \frac{40000}{20} \right\rceil = 2000$$

and the index holds block pointers for each of these 2000 blocks. The number of index block accesses is

$$1 + \left\lceil \frac{2000}{200} \right\rceil = 11$$

Therefore, the minimum number of block accesses is = $2000 + 11 = 2011$

In the worst case, we access at most 40000 blocks³. 4K block pointers

³We cannot partition 40000 rows into 50000 blocks such that each block gets at least one row

are stored in the index. The number of index blocks accessed is

$$1 + \left\lceil \frac{40000}{200} \right\rceil = 201$$

Therefore, the maximum number of block accesses is = $40000 + 201 = 40201$

YES1 = 1

Both in the best and worst case, one data block and two index blocks are accessed.

Therefore, the maximum and minimum number of block accesses is = $1 + 2 = 3$

# YES	Minimum	Maximum
999981	50250	1004981
40001	2011	40201
1	3	3

There is another argument to be made for the maximum number of block accesses. The textbook discusses an implementation of the secondary index that involves storing $\langle K, P \rangle$ pairs explicitly where the key K is repeated for each pointer. In such a case, each block can only store 100 pointers and not 200 as assumed before. If this argument is made, I present the results below

# YES	Minimum	Maximum
999981	50500	1009981
40001	2021	40401
1	3	3

2022 DATA SYSTEMS FINAL EXAM (SOLUTIONS) 6

Question 3

Consider four relations R_1, R_2, R_3, R_4 like in **Question 1**, with attribute A in all relations and domain values of A as $dom(A) = \{1, 2, 3, 4, 5, 6, 7\}$. Let n_{ij} values be:

$dom(A)$	R_1	R_2	R_3	R_4
1	10	10	40	50
2	20	20	80	60
3	20	40	90	40
4	100	0	100	70
5	10	20	20	80
6	10	10	30	50
7	20	30	40	30

Consider the following query:

```
SELECT R1.A
FROM R1, R2, R3, R4
WHERE R1.A = R2.A AND R2.A = R3.A AND R3.A = R4.A
AND R1.A > 3 AND R2.A > 2 AND R3.A < 5 AND R4.A = 4
```

Come up with the best possible cost-based query optimization plan with sizes of the intermediate results (in the number of rows). Give at most three sentences of justification for your query execution plan. What is the result of the query?

Answer. After using the heuristics detailed in Fig. 700 of the textbook⁴, the optimization problem boils down to a join ordering problem as shown below. The intermediate sizes are shown post the selection and projection operators. Since the histogram is provided, I assume the query optimizer uses the histogram values to estimate cost.

```

graph TD
    S1[140] --> S2[130]
    S1 --> S3[130]
    S2 --> S4[130]
    S2 --> S5[130]
    S3 --> S6[130]
    S3 --> S7[130]
    S4 --> S8[130]
    S4 --> S9[130]
    S5 --> S10[130]
    S5 --> S11[130]
    S6 --> S12[130]
    S6 --> S13[130]
    S7 --> S14[130]
    S7 --> S15[130]
    S8 --> S16[130]
    S8 --> S17[130]
    S9 --> S18[130]
    S9 --> S19[130]
    S10 --> S20[130]
    S10 --> S21[130]
    S11 --> S22[130]
    S11 --> S23[130]
    S12 --> S24[130]
    S12 --> S25[130]
    S13 --> S26[130]
    S13 --> S27[130]
    S14 --> S28[130]
    S14 --> S29[130]
    S15 --> S30[130]
    S15 --> S31[130]
    S16 --> S32[130]
    S16 --> S33[130]
    S17 --> S34[130]
    S17 --> S35[130]
    S18 --> S36[130]
    S18 --> S37[130]
    S19 --> S38[130]
    S19 --> S39[130]
    S20 --> S40[130]
    S20 --> S41[130]
    S21 --> S42[130]
    S21 --> S43[130]
    S22 --> S44[130]
    S22 --> S45[130]
    S23 --> S46[130]
    S23 --> S47[130]
    S24 --> S48[130]
    S24 --> S49[130]
    S25 --> S50[130]
    S25 --> S51[130]
    S26 --> S52[130]
    S26 --> S53[130]
    S27 --> S54[130]
    S27 --> S55[130]
    S28 --> S56[130]
    S28 --> S57[130]
    S29 --> S58[130]
    S29 --> S59[130]
    S30 --> S60[130]
    S30 --> S61[130]
    S31 --> S62[130]
    S31 --> S63[130]
    S32 --> S64[130]
    S32 --> S65[130]
    S33 --> S66[130]
    S33 --> S67[130]
    S34 --> S68[130]
    S34 --> S69[130]
    S35 --> S70[130]
    S35 --> S71[130]
    S36 --> S72[130]
    S36 --> S73[130]
    S37 --> S74[130]
    S37 --> S75[130]
    S38 --> S76[130]
    S38 --> S77[130]
    S39 --> S78[130]
    S39 --> S79[130]
    S40 --> S80[130]
    S40 --> S81[130]
    S41 --> S82[130]
    S41 --> S83[130]
    S42 --> S84[130]
    S42 --> S85[130]
    S43 --> S86[130]
    S43 --> S87[130]
    S44 --> S88[130]
    S44 --> S89[130]
    S45 --> S90[130]
    S45 --> S91[130]
    S46 --> S92[130]
    S46 --> S93[130]
    S47 --> S94[130]
    S47 --> S95[130]
    S48 --> S96[130]
    S48 --> S97[130]
    S49 --> S98[130]
    S49 --> S99[130]
    S50 --> S100[130]
    S50 --> S101[130]
    S51 --> S102[130]
    S51 --> S103[130]
    S52 --> S104[130]
    S52 --> S105[130]
    S53 --> S106[130]
    S53 --> S107[130]
    S54 --> S108[130]
    S54 --> S109[130]
    S55 --> S110[130]
    S55 --> S111[130]
    S56 --> S112[130]
    S56 --> S113[130]
    S57 --> S114[130]
    S57 --> S115[130]
    S58 --> S116[130]
    S58 --> S117[130]
    S59 --> S118[130]
    S59 --> S119[130]
    S60 --> S120[130]
    S60 --> S121[130]
    S61 --> S122[130]
    S61 --> S123[130]
    S62 --> S124[130]
    S62 --> S125[130]
    S63 --> S126[130]
    S63 --> S127[130]
    S64 --> S128[130]
    S64 --> S129[130]
    S65 --> S130[130]
    S65 --> S131[130]
    S66 --> S132[130]
    S66 --> S133[130]
    S67 --> S134[130]
    S67 --> S135[130]
    S68 --> S136[130]
    S68 --> S137[130]
    S69 --> S138[130]
    S69 --> S139[130]
    S70 --> S140[130]
    S70 --> S141[130]
    S71 --> S142[130]
    S71 --> S143[130]
    S72 --> S144[130]
    S72 --> S145[130]
    S73 --> S146[130]
    S73 --> S147[130]
    S74 --> S148[130]
    S74 --> S149[130]
    S75 --> S150[130]
    S75 --> S151[130]
    S76 --> S152[130]
    S76 --> S153[130]
    S77 --> S154[130]
    S77 --> S155[130]
    S78 --> S156[130]
    S78 --> S157[130]
    S79 --> S158[130]
    S79 --> S159[130]
    S80 --> S160[130]
    S80 --> S161[130]
    S81 --> S162[130]
    S81 --> S163[130]
    S82 --> S164[130]
    S82 --> S165[130]
    S83 --> S166[130]
    S83 --> S167[130]
    S84 --> S168[130]
    S84 --> S169[130]
    S85 --> S170[130]
    S85 --> S171[130]
    S86 --> S172[130]
    S86 --> S173[130]
    S87 --> S174[130]
    S87 --> S175[130]
    S88 --> S176[130]
    S88 --> S177[130]
    S89 --> S178[130]
    S89 --> S179[130]
    S90 --> S180[130]
    S90 --> S181[130]
    S91 --> S182[130]
    S91 --> S183[130]
    S92 --> S184[130]
    S92 --> S185[130]
    S93 --> S186[130]
    S93 --> S187[130]
    S94 --> S188[130]
    S94 --> S189[130]
    S95 --> S190[130]
    S95 --> S191[130]
    S96 --> S192[130]
    S96 --> S193[130]
    S97 --> S194[130]
    S97 --> S195[130]
    S98 --> S196[130]
    S98 --> S197[130]
    S99 --> S198[130]
    S99 --> S199[130]
    S100 --> S200[130]
    S100 --> S201[130]
    S101 --> S202[130]
    S101 --> S203[130]
    S102 --> S204[130]
    S102 --> S205[130]
    S103 --> S206[130]
    S103 --> S207[130]
    S104 --> S208[130]
    S104 --> S209[130]
    S105 --> S210[130]
    S105 --> S211[130]
    S106 --> S212[130]
    S106 --> S213[130]
    S107 --> S214[130]
    S107 --> S215[130]
    S108 --> S216[130]
    S108 --> S217[130]
    S109 --> S218[130]
    S109 --> S219[130]
    S110 --> S220[130]
    S110 --> S221[130]
    S111 --> S222[130]
    S111 --> S223[130]
    S112 --> S224[130]
    S112 --> S225[130]
    S113 --> S226[130]
    S113 --> S227[130]
    S114 --> S228[130]
    S114 --> S229[130]
    S115 --> S230[130]
    S115 --> S231[130]
    S116 --> S232[130]
    S116 --> S233[130]
    S117 --> S234[130]
    S117 --> S235[130]
    S118 --> S236[130]
    S118 --> S237[130]
    S119 --> S238[130]
    S119 --> S239[130]
    S120 --> S240[130]
    S120 --> S241[130]
    S121 --> S242[130]
    S121 --> S243[130]
    S122 --> S244[130]
    S122 --> S245[130]
    S123 --> S246[130]
    S123 --> S247[130]
    S124 --> S248[130]
    S124 --> S249[130]
    S125 --> S250[130]
    S125 --> S251[130]
    S126 --> S252[130]
    S126 --> S253[130]
    S127 --> S254[130]
    S127 --> S255[130]
    S128 --> S256[130]
    S128 --> S257[130]
    S129 --> S258[130]
    S129 --> S259[130]
    S130 --> S260[130]
    S130 --> S261[130]
    S131 --> S262[130]
    S131 --> S263[130]
    S132 --> S264[130]
    S132 --> S265[130]
    S133 --> S266[130]
    S133 --> S267[130]
    S134 --> S268[130]
    S134 --> S269[130]
    S135 --> S270[130]
    S135 --> S271[130]
    S136 --> S272[130]
    S136 --> S273[130]
    S137 --> S274[130]
    S137 --> S275[130]
    S138 --> S276[130]
    S138 --> S277[130]
    S139 --> S278[130]
    S139 --> S279[130]
    S140 --> S280[130]
    S140 --> S281[130]
    S141 --> S282[130]
    S141 --> S283[130]
    S142 --> S284[130]
    S142 --> S285[130]
    S143 --> S286[130]
    S143 --> S287[130]
    S144 --> S288[130]
    S144 --> S289[130]
    S145 --> S290[130]
    S145 --> S291[130]
    S146 --> S292[130]
    S146 --> S293[130]
    S147 --> S294[130]
    S147 --> S295[130]
    S148 --> S296[130]
    S148 --> S297[130]
    S149 --> S298[130]
    S149 --> S299[130]
    S150 --> S300[130]
    S150 --> S301[130]
    S151 --> S302[130]
    S151 --> S303[130]
    S152 --> S304[130]
    S152 --> S305[130]
    S153 --> S306[130]
    S153 --> S307[130]
    S154 --> S308[130]
    S154 --> S309[130]
    S155 --> S310[130]
    S155 --> S311[130]
    S156 --> S312[130]
    S156 --> S313[130]
    S157 --> S314[130]
    S157 --> S315[130]
    S158 --> S316[130]
    S158 --> S317[130]
    S159 --> S318[130]
    S159 --> S319[130]
    S160 --> S320[130]
    S160 --> S321[130]
    S161 --> S322[130]
    S161 --> S323[130]
    S162 --> S324[130]
    S162 --> S325[130]
    S163 --> S326[130]
    S163 --> S327[130]
    S164 --> S328[130]
    S164 --> S329[130]
    S165 --> S330[130]
    S165 --> S331[130]
    S166 --> S332[130]
    S166 --> S333[130]
    S167 --> S334[130]
    S167 --> S335[130]
    S168 --> S336[130]
    S168 --> S337[130]
    S169 --> S338[130]
    S169 --> S339[130]
    S170 --> S340[130]
    S170 --> S341[130]
    S171 --> S342[130]
    S171 --> S343[130]
    S172 --> S344[130]
    S172 --> S345[130]
    S173 --> S346[130]
    S173 --> S347[130]
    S174 --> S348[130]
    S174 --> S349[130]
    S175 --> S350[130]
    S175 --> S351[130]
    S176 --> S352[130]
    S176 --> S353[130]
    S177 --> S354[130]
    S177 --> S355[130]
    S178 --> S356[130]
    S178 --> S357[130]
    S179 --> S358[130]
    S179 --> S359[130]
    S180 --> S360[130]
    S180 --> S361[130]
    S181 --> S362[130]
    S181 --> S363[130]
    S182 --> S364[130]
    S182 --> S365[130]
    S183 --> S366[130]
    S183 --> S367[130]
    S184 --> S368[130]
    S184 --> S369[130]
    S185 --> S370[130]
    S185 --> S371[130]
    S186 --> S372[130]
    S186 --> S373[130]
    S187 --> S374[130]
    S187 --> S375[130]
    S
```